



Introduction to Virtual Machines

[1] E. Bugnion, J. Nieh, and D. Tsafrir, “[Hardware and Software Support For Virtualization](#),” *Synth. Lect. Comput. Archit.*, vol. 12, no. 1, pp. 1–206, Feb. 2017.

Readings: Chapter 1 [1]

Abstraction and Layering

- ▣ **Abstraction**: only way of dealing with complex systems
 - ◆ Divide world into objects, each with an...
 - **Interface**: knobs, behaviors, knobs → behaviors
 - **Implementation**: “black box”
 - ◆ Specialists deal with implementation; others interface
 - ◆ Example: car drivers vs. mechanics

- ▣ **Layering**: abstraction discipline makes life even simpler
 - ◆ Removes need to even know interfaces of most objects
 - ◆ Divide objects in system into layers
 - ◆ Layer X objects
 - Implemented in terms of interfaces of layer X-1 objects
 - Don't even need to know interfaces of layer X-2 objects
 - ◆ Example: cab passenger vs. mechanics

Abstraction, Layering, and Computers

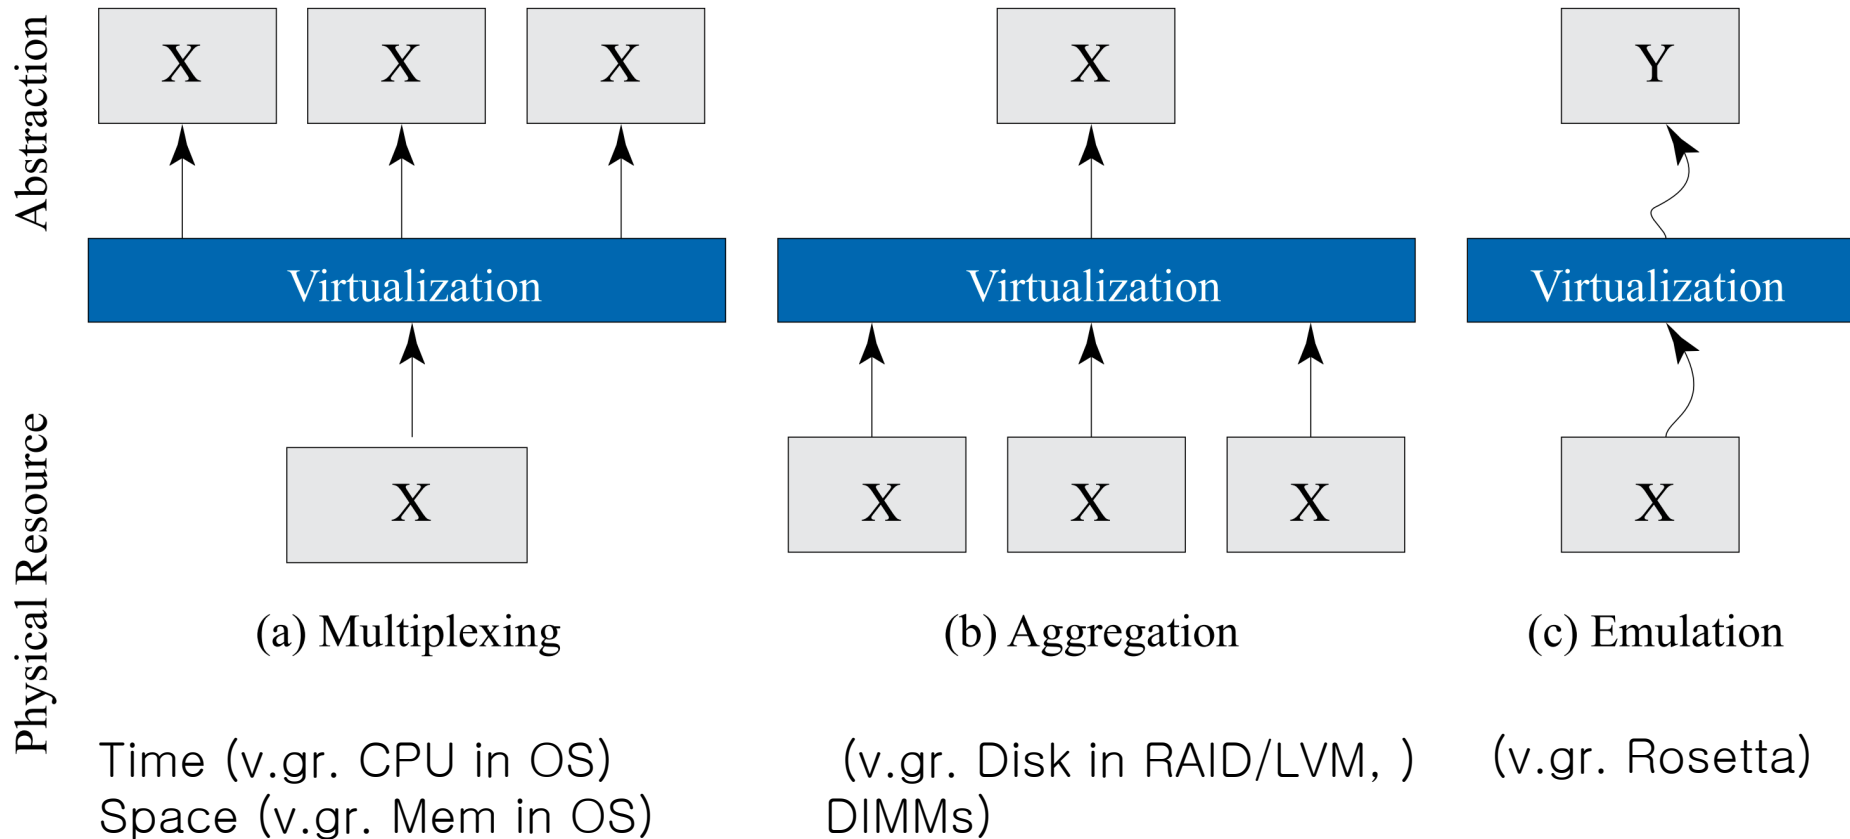
- ▣ Computers are complex systems, built in layers
 - ◆ Applications
 - ◆ O/S, compiler
 - ◆ Firmware, device drivers
 - ◆ Processor, memory, raw I/O devices
 - ◆ Digital circuits, digital/analog converters
 - ◆ Gates
 - ◆ Transistors
- ▣ 99% of users don't know hardware layers implementation
- ▣ 90% of users don't know implementation of any layer
- ▣ That's OK, world still works just fine
 - ◆ But unfortunately, the layers sometimes breakdown
 - ◆ Someone needs to understand what's "under the hood"

Virtualization

- ▣ Definition:
 - ◆ *Virtualization is the application of the layering principle through **enforced modularity**, whereby the exposed virtual resource is identical to the underlying physical resource being virtualized.*
- ▣ Virtualization Example in Computer Architecture
 - ◆ Virtual memory modularity enforced through MMU
- ▣ Virtualization within the Operating System
 - ◆ Expose real resources (CPU, Memory, I/O) to processes in a controlled way
- ▣ Virtualization in I/O subsystems
 - ◆ RAID controllers and Disks

Implementation Techniques in Virtualization

- Combination of these three techniques in the hypervisor

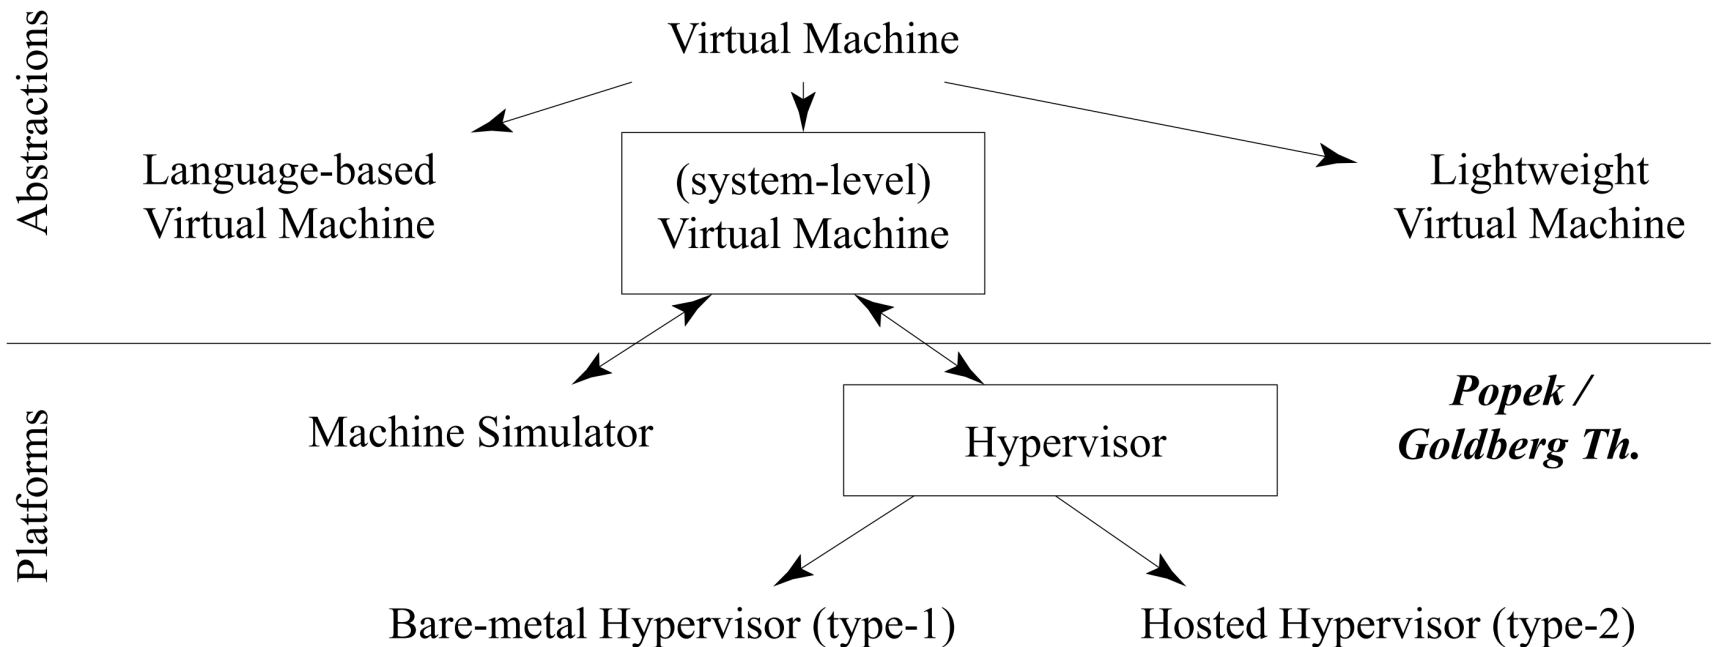


Virtual Machines: Definitions

- ▣ **Virtualization** is the application of the layering principle through enforced modularity, whereby the exposed virtual resource is identical to the underlying physical resource being virtualized
- ▣ **A virtual machine** is an abstraction of a complete compute environment through the combined virtualization of the processor, memory, and I/O components of a computer.
- ▣ **The hypervisor** is a specialized piece of system software that manages and runs virtual machines.
- ▣ **The virtual machine monitor** (VMM) refers to the portion of the hypervisor that focuses on the CPU and memory virtualization (beware)

Virtual Machines

- A **virtual machine** is a complete compute environment with its own isolated processing capabilities, memory, and communication channels.



Classes of Virtual Machines

- ▣ Language-based virtual machines
 - ◆ For portability reasons, synthetic ISA to compile.
 - ◆ ISA is translated on runtime into hardware ISA
 - ◆ Not our focus
 - ◆ Ej: JVM

- ▣ Lightweight virtual machines,
 - ◆ Rely on a combination of hardware and software isolation mechanisms to ensure that applications **running directly on the processor** (e.g., as native x86 code) are securely isolated.
 - ◆ Ej: containers

- ▣ **System Level Virtual Machines**

System Level Virtual Machines

- Compute environment that **resembles** the hardware of a computer with enough detail to run a standard, **commodity operating** system and its applications
- **Full isolation** from the other virtual machines and the rest of the environment.
- Applies the virtualization principle to **an entire computer system**.
- Each virtual machine has its own copy of the underlying hardware, or at least, its own copy of **some**
- Each virtual machine runs its own independent operating system instance, called the **guest operating system**

Platforms

- ▣ Machine Simulators (full system simulators)
 - ◆ Implemented as a user-level application
 - ◆ Models functionally (and timing) hardware details of the platform to study (processor, memory, I/O, etc). Platform can be single-system or multi-system
 - ◆ Very slow
 - ◆ Useful to hardware/software codesign, computer architecture research, etc..
 - ◆ Example: **gem5**

- ▣ Hypervisor
 - ◆ Relies on direct execution on the CPU
 - ◆ Should emulate non-user level instructions somehow
 - ◆ Models functionally other less performance sensitive components (Disk, net)
 - ◆ ... if supported, can be used directly
 - ◆ Example: **Xen, KVM**

Hypervisor

▣ Popek and Goldberg Definition

A virtual machine is taken to be an **efficient, isolated duplicate** of the real machine. We explain these notions through the idea of a virtual machine monitor (VMM). As a piece of software, a VMM has three essential characteristics. First, the VMM provides an environment for programs which is essentially identical with the original machine; second, programs running in this environment show at worst only minor decreases in speed; and last, the VMM is in complete control of system resources.

▣ Equivalence

- ◆ Duplicating real resources

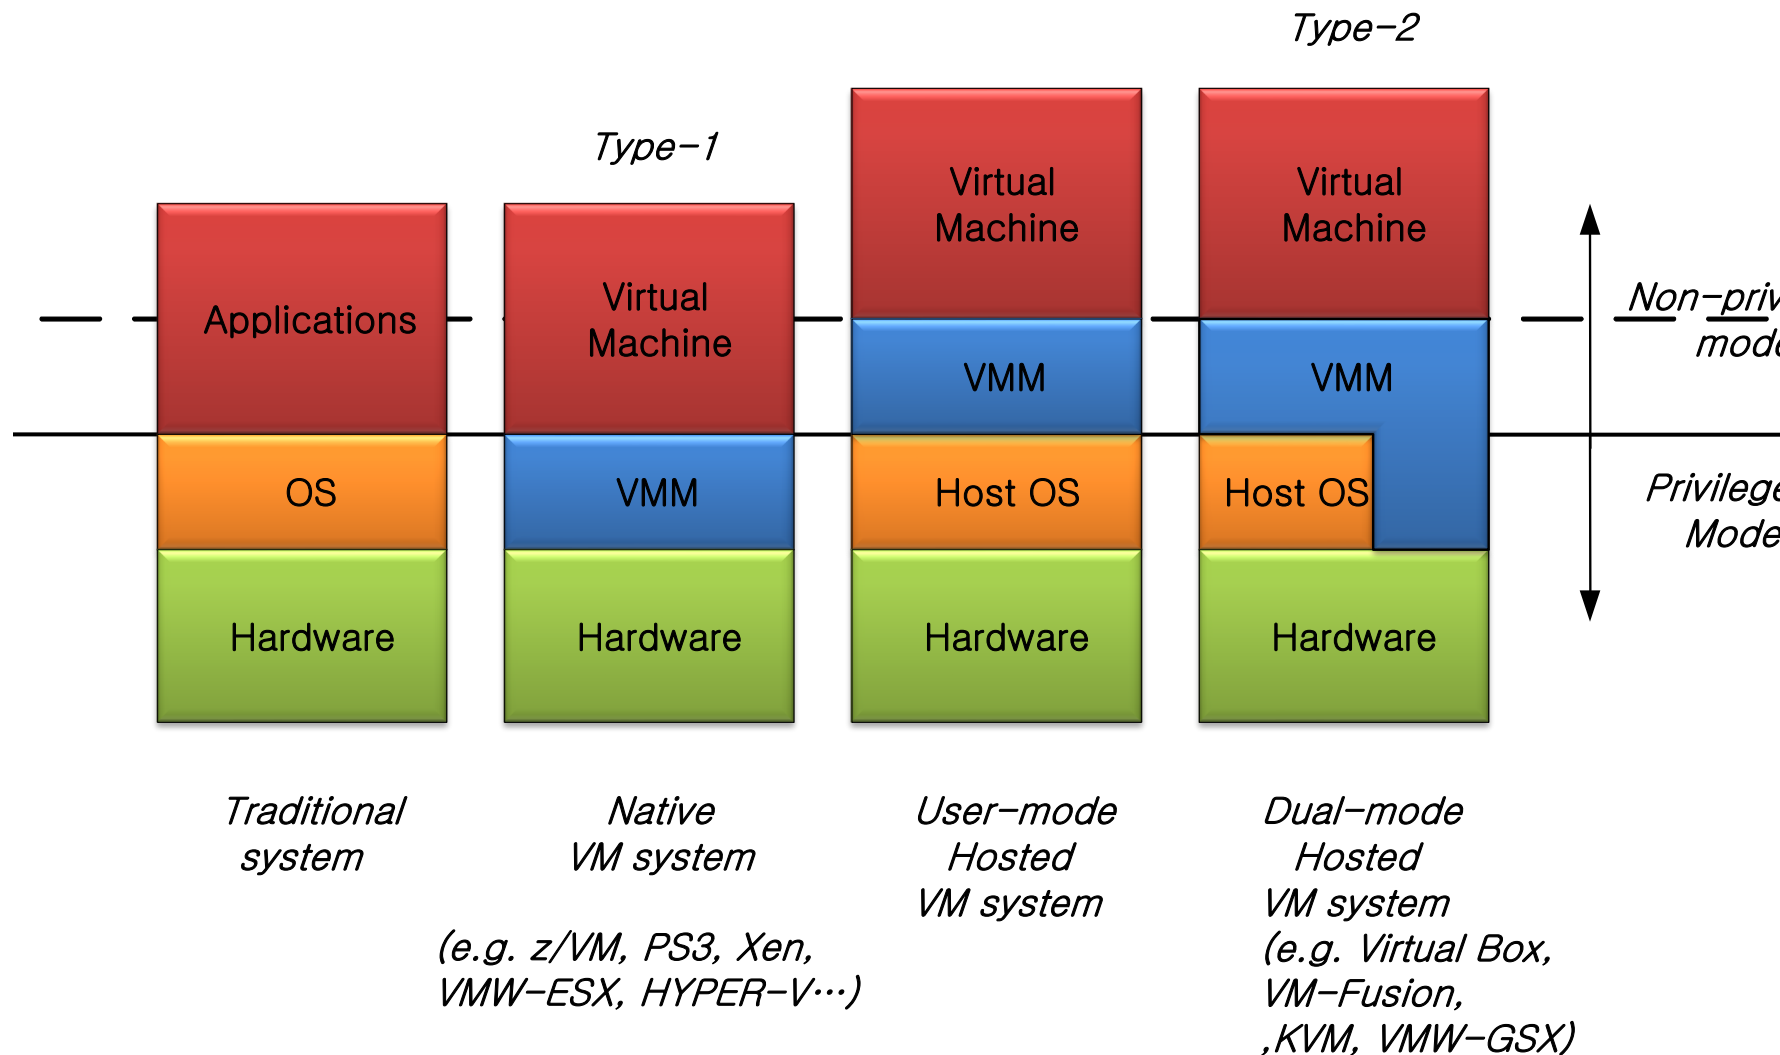
▣ Safety

- ◆ Isolation between VM and hypervisor

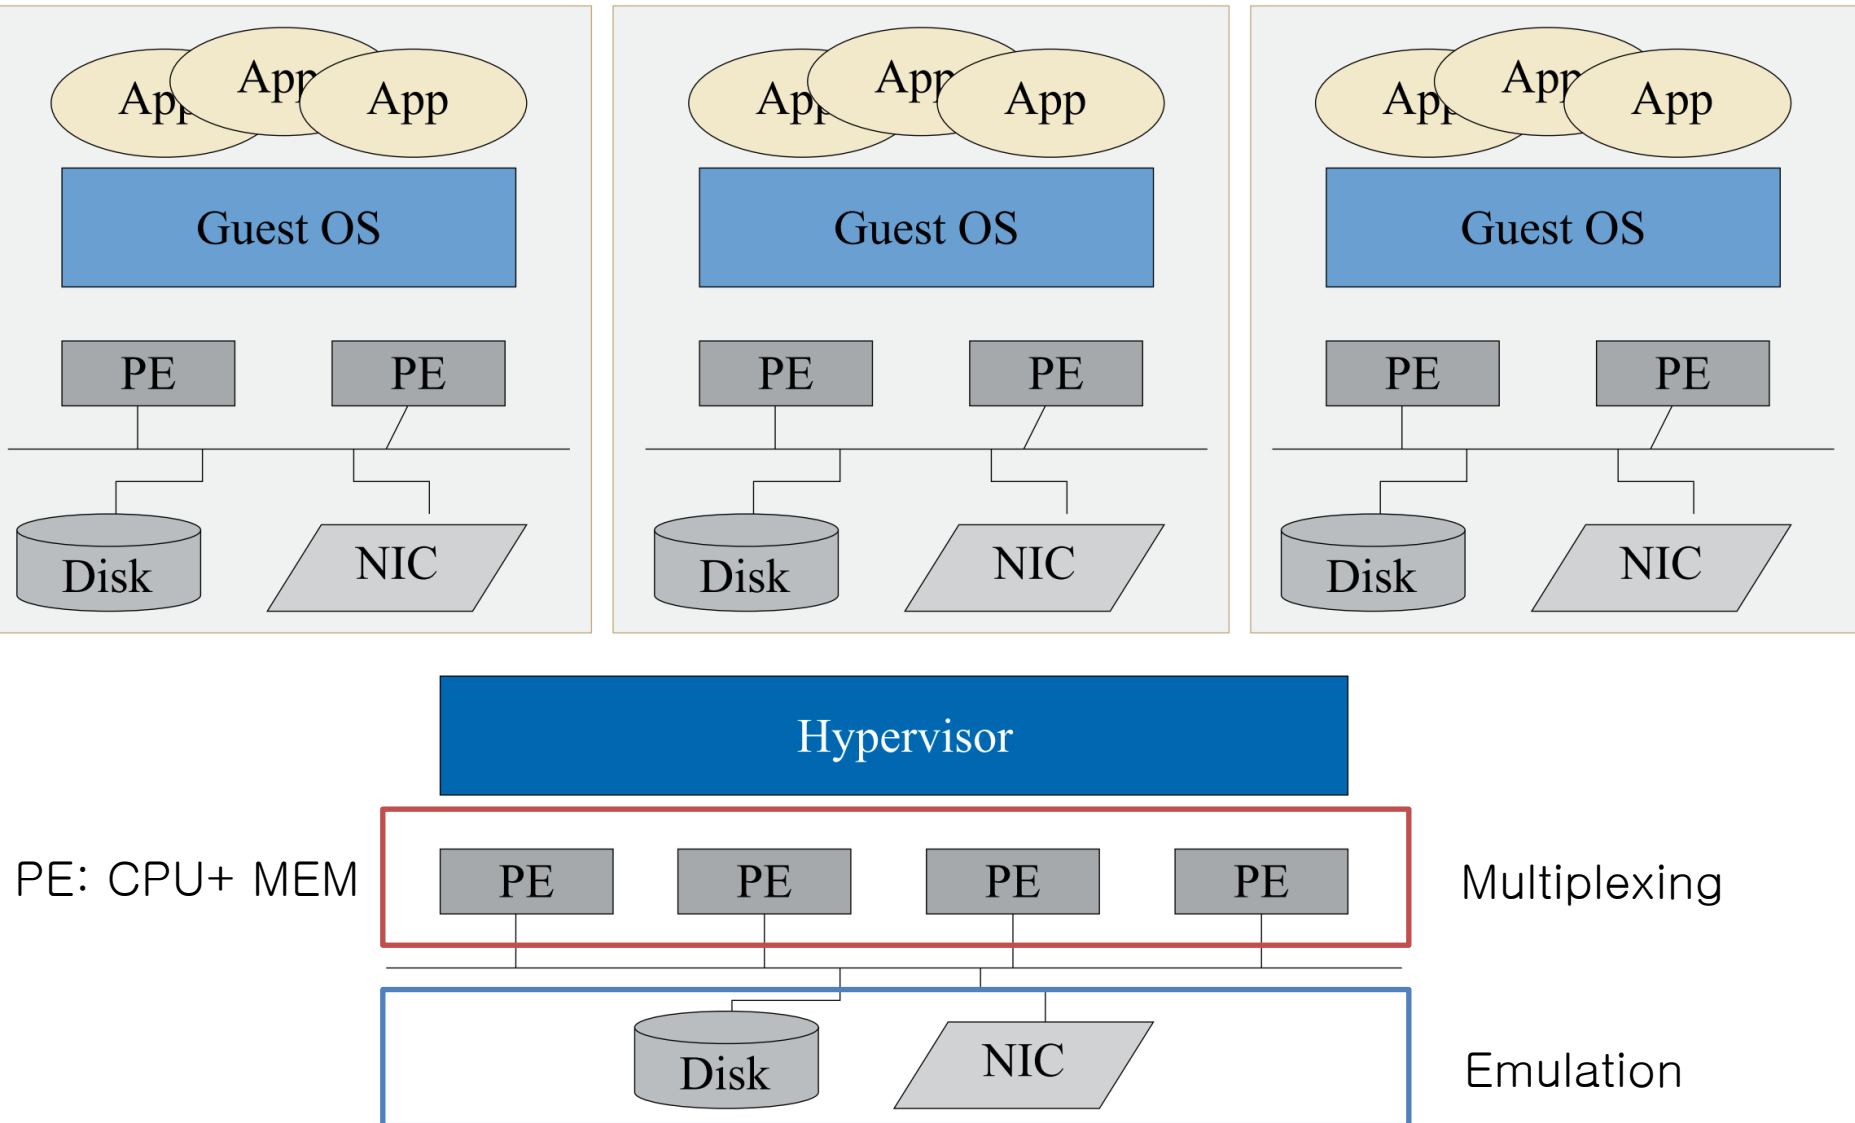
▣ Performance

- ◆ Separates Hypervisors from Simulators

Type-1 and Type-2 hypervisors



Multiplexing and Emulation



Case of VMWare (Early versions, No HW support)

	Virtual Hardware (front-end)	Back-end
Multiplexed	1 virtual x86-32 CPU	Scheduled by the host operating system with one or more x86 CPUs
	Up to 512 MB of contiguous DRAM	Allocated and managed by the host OS (page-by-page)
Emulated	PCI Bus	Fully emulated compliant PCI bus with B/D/F addressing for all virtual motherboard and slot devices
	4 x 4IDE disks 7 x Buslogic SCSI Disks	Either virtual disks (stored as files) or direct access to a given raw device
	1 x IDE CD-ROM	ISO image or real CD-ROM
	2 x 1.44 MB floppy drives	Physical floppy or floppy image
	1 x VGA/SVGA graphics card	Appears as a Window or in full-screen mode
	2 x serial ports COM1 and COM2	Connect to Host serial port or a file
	1 x printer (LPT)	Can connect to host LPT port
	1 x keyboard (104-key) and mouse	Fully emulated
	AMD PCnet NIC (AM79C970A)	Via virtual switch of the host

Names for Memory

- Eskimos and Snow, Computer Architects and Memory
 - ◆ **Virtual memory** concept is the most significant enhancement over the original Von-Neumman Model

- Virtual memory
 - ◆ Byte addressable **namespace** used by instruction sequences executed by the processor

- Physical Memory
 - ◆ Byte addressable resource accessed by the memory hierarchy (typically DRAM)
 - ◆ **Guest-physical** memory or **Host-physical** Memory in a VM

Approaches to Virtualization

▣ Full (software) **virtualization**

- ◆ Hypervisors designed to maximize hardware compatibility
- ◆ Run unmodified operating systems on architectures without full support for it (usually by the means of dynamic binary translation)

▣ **Hardware Virtualization** (HVM)

- ◆ Hypervisors built for hardware with architectural support for virtualization
- ◆ (mostly) Rely on direct execution

▣ **Paravirtualization**

- ◆ Initially, hypervisors for platform without architectural support for virtualization using modified operating system to avoid binary translation (via hyper-calls)
- ◆ Today, a mix of paravirtualization and HVM is usual

Benefits of Virtual Machines

- ▣ Operating system diversity (on same hardware)
- ▣ Server consolidation
 - ◆ Best IT practices mandates single app per “server”
- ▣ Rapid provisioning
 - ◆ Simplify server deployment
- ▣ Security
 - ◆ Additional OS Isolation, OS introspection, external control (e.g. firewall interposing)
- ▣ High-availability
 - ◆ Near-zero operation impact of hardware disruption
- ▣ Distributed resource scheduling
 - ◆ Live migration techniques convert a group of hypervisors in a common resource pool
- ▣ Cloud computing
 - ◆ Mix customers (tenants) in a shared resource pool. Requires **network virtualization**

History of sVM

- ▣ Mid/late-60s IBM 360/67 -- CP-67
 - ◆ First 360 with VM.
 - ◆ CMS an essential part (Multiuser in a single user OS)
- ▣ Late 60s/early 70s
 - ◆ VMs blossomed as a research topic
- ▣ Early 70s several VM implementations
 - ◆ Honeywell
 - ◆ DEC
 - ◆ RCA
 - ◆ Several university projects
- ▣

sVM on cheap HW (x86)

▣ Software Explosion

- ◆ VMWare (v1.0 1999)
 - Targets Windows
 - From Stanford Disco Project (mainly SimOS)
- ◆ Xen (v1.0 2003)
 - Targets Linux
 - AWS
 - From Cambridge Univ.

▣ Hardware

- ◆ Intel: Vanderpool (VT-x): Pentium 4 (2005)
- ◆ AMD: Pacifica (AMD-v) Athlon64 Processors (2006)