

# X86-64: MMU Virtualization With EPT

---

[Hardware and Software Support For Virtualization](#)

Chapter 5

# Memory is a Two-dimensional problem

## ▣ Problem

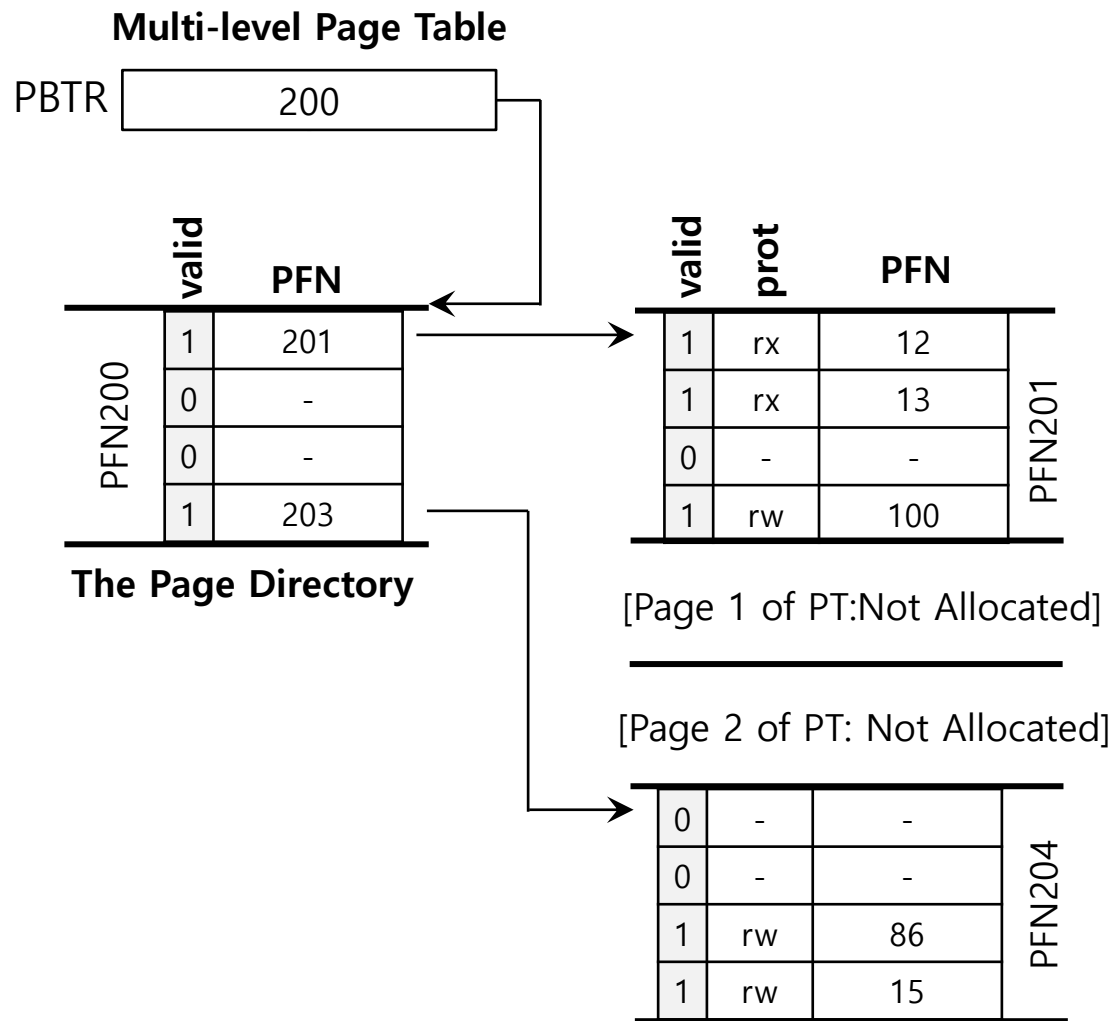
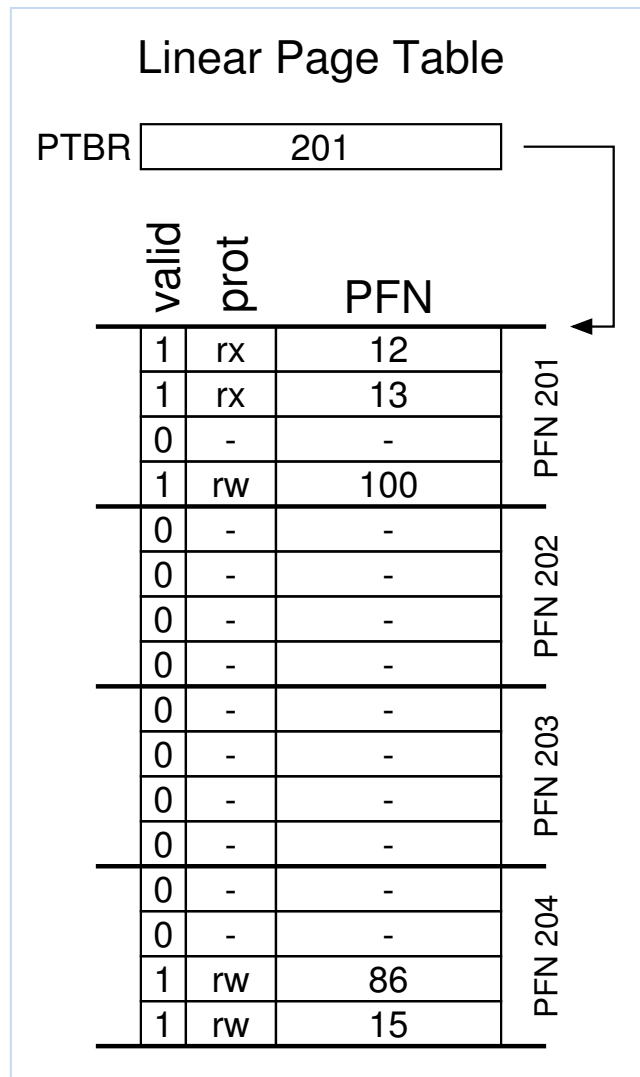
- ◆ Each VM is provided with the **guest-physical** memory abstraction
- ◆ Hypervisor handles **host-physical** memory (i.e., the physical resource)

## ▣ Memory translation is no longer a linear problem: a unidimensional approach (i.e., *linear* page table) **won't suffice**

## ▣ Approaches

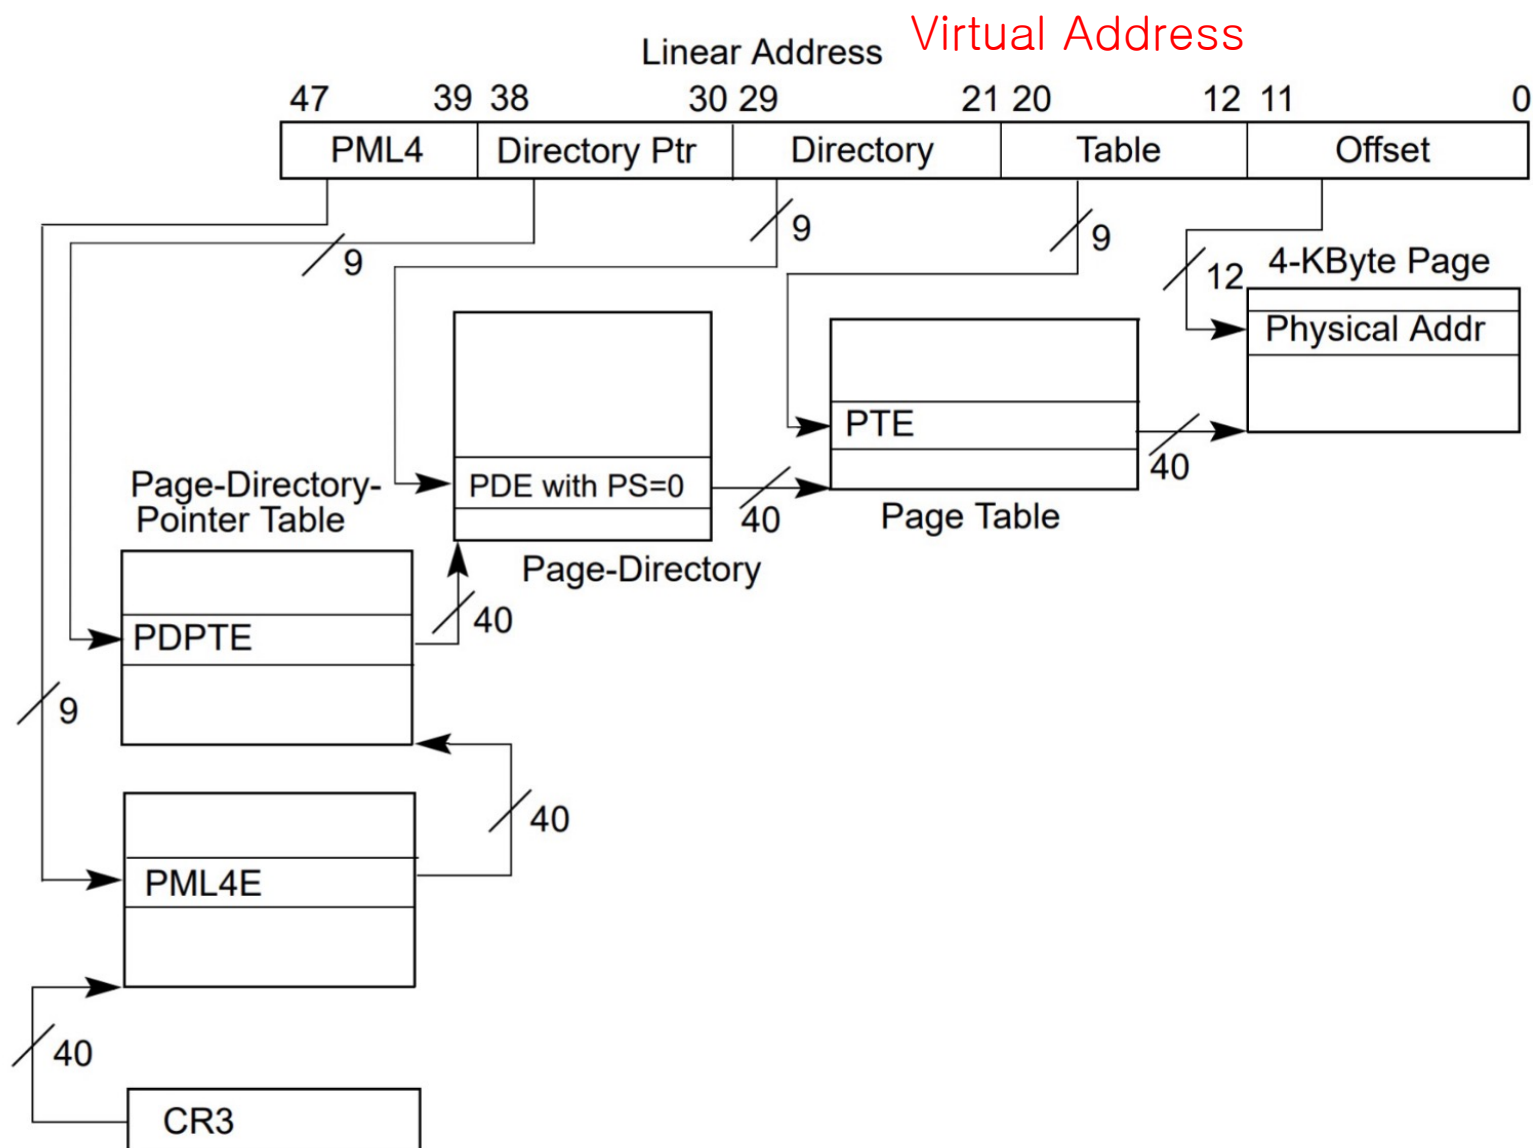
- ◆ **Pure Software solutions** on HW conceived for two privilege levels (OS) are required (e.g. Shadow Page Tables). Arguably **the most complex** part of the Hypervisor. [ Memory paravirtualization restricted by guest-OS/VMM ]
- ◆ **Hardware support** for bi-dimensional page tables, (a.k.a. Extended Page Tables or Nested Page Tables)

# Review: Multi Level Page Tables I




**Linear (Left) And Multi-Level (Right) Page Tables**


# Review: Multilevel Page Tables II ( $2^{48}B=256TB$ )




# Mega pages (2MB pages)



# Giga Pages (1GB Pages)




# It Matters?



[1] A. Basu, J. Gandhi, J. Chang, M. D. Hill, and M. M. Swift, “Efficient Virtual Memory for Big Memory Servers”, CAL 2013.

# Review: Since 2021 ( $2^{56}\text{B}=128\text{PB}$ )

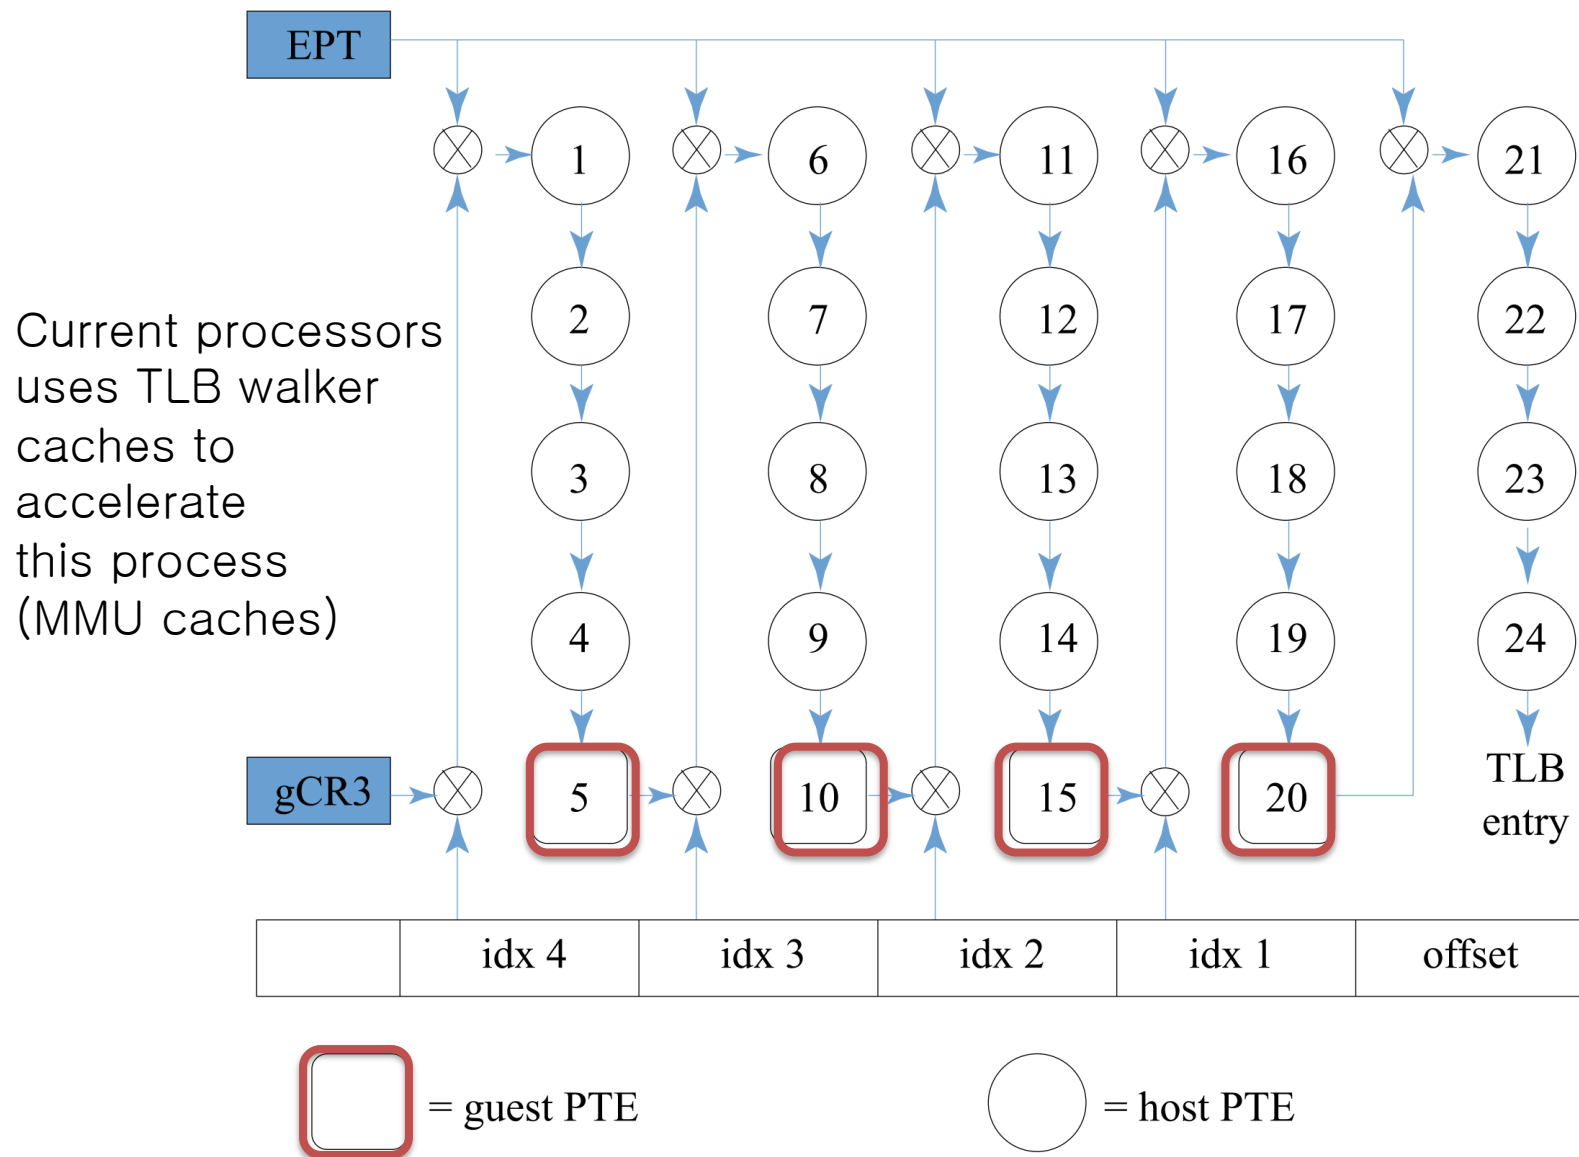




# Extended Page Tables

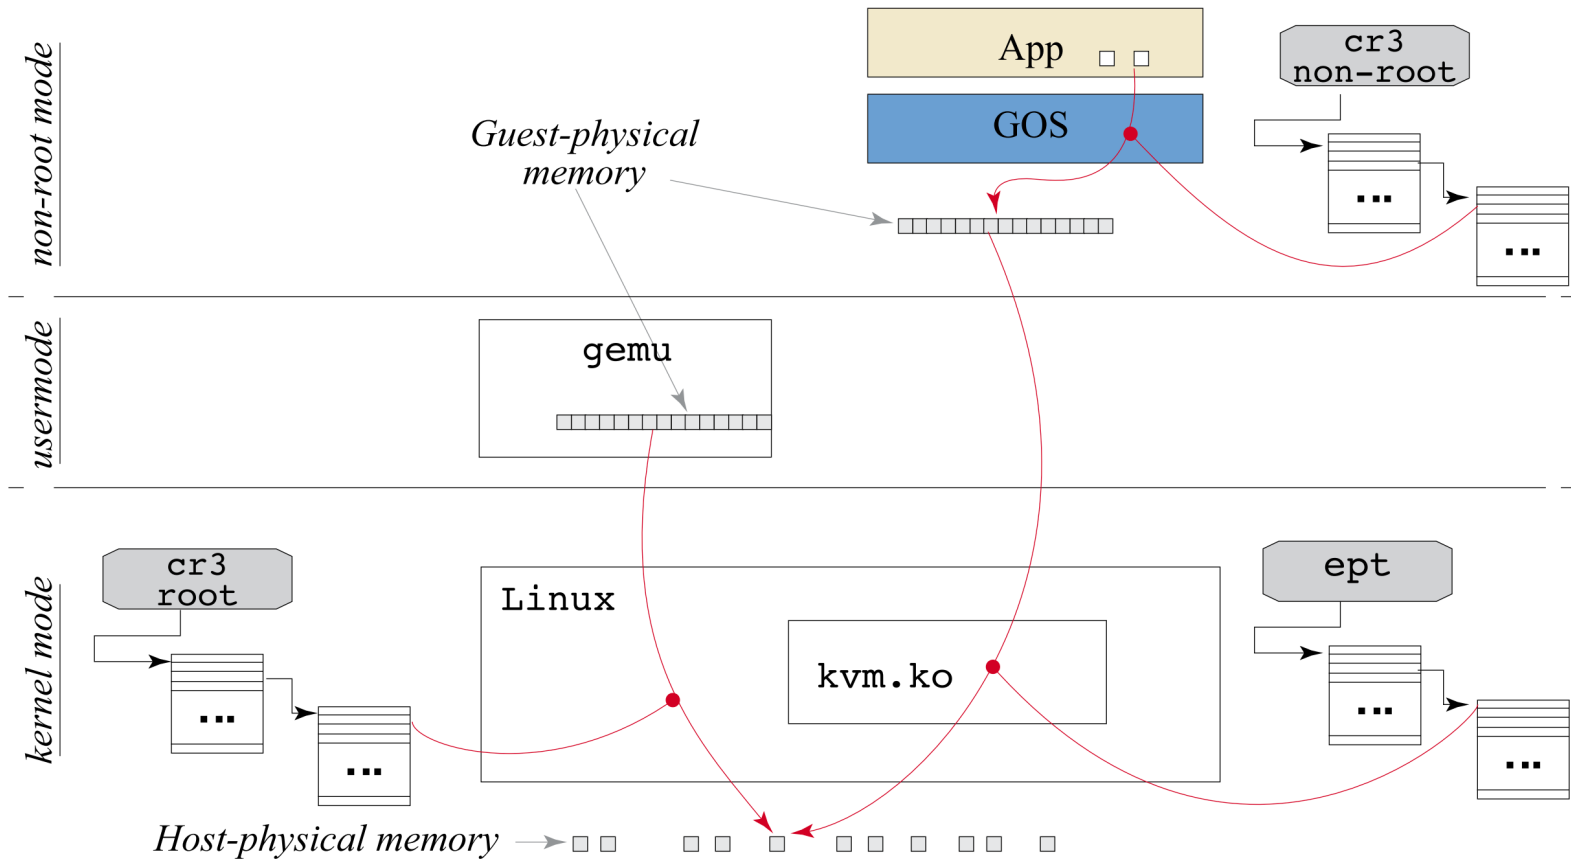
- Circa 2008 Intel (Extended Page Tables or **EPT**) and AMD (Nested Page Tables or **NPT**)
- Add hardware support to complement the classic hardware-defined page table structures (maintained guest-OS) with a **second page table** structure (maintained by the **Hypervisor**)
- TLB miss handling logic is deeply integrated with it (i.e., when we assign the `%cr3`, TLB faults are resolved by TLB walker under the table)
- Memory accesses after TLB miss
  - ◆ In root mode EPT is disabled tree levels  $n=4$  (kB pages),  $n=3$ , (2MB pages),  $n=2$  (1GB pages), will be  $n$  access
  - ◆ In non-root mode, another table (host-to-physical) with  $m$  levels will be  $n \times m + 1 + m$

# Example: m=4, n=4 (4KB EPT Pages, 4KB Mem Page)



# EPT in KVM

- KVM was conceived when EPT wasn't available: now the feature can be used or not. Today non using it is mostly restricted to nested virtualization




# Actors in the memory access pass

- QEMU: Allocates guest-physical **contiguous memory** in his own virtual address space
  - ◆ The details about memory management is deferred to the host-OS (type-2)
  - ◆ Provides a convenient access from user space to the guest-physical space (crucial for DMA device emulation)
  - ◆ (as any other process) will have a `%cr3` in the root mode
- Virtual machine manages its own page tables. Non-root `%cr3` points to the page table of the process running (1cpu)
  - ◆ Assignments (e.g. context switch) of non-root `%cr3` do not cause `#vmexit`
- In root mode there is special register (`eptp`) pointing to the EPT.  
`kvm.ko` **manages** the content of it (one per VM, i.e., accessible via VMCS)
  - ◆ One `eptp` per VM

# Collateral complications


- ▣ **%eip** in App uses virtual addresses: hypervisor needs to decode the mapping before emulate the instruction
  - ◆ It's trivial to access to the guest-physical address space (add a constant)
  - ◆ **Its very hard to access** (to emulate) the virtual address space of the App (translations are only at the non-root MMU but non on root [**hypervisor**])
- ▣ Type-2 KVM design required a constant synchronization between QEMU and guest-OS mapping
  - ◆ E.g. If a page in QEMU is swapped out, affected guest-OS page table should be updated
- ▣ **KSM** mechanism (allows to **share read-only pages** within processes) can be used (i.e multiple VM can share host physical pages)



# Performance Considerations

- EPT avoids 80% `vmexit`, but TLB misses are much more costly
  - ◆ In 64-bit address space, with 4KB ( $m=n=4$ ), a single TLB miss requires 24 access to memory
  - ◆ Page walker caches reduce it (hardware)
- In spite of the efforts, some pointer-chasing benchmarks showed 17% and 39% **performance degradation** on Intel and AMD in early designs
- Next, similar pointer-chasing with Intel **Sandy Bridge** Architectures
  - ◆ List of objects that randomly links to a heap of 32GB
  - ◆ Measure average access time to each object

# EPT Configuration Effect on Sandy Bridge (non NUMA effects)



# Tradeoffs in EPT

- ▣ Although prev. benchmark is somewhat unrealistic, **TLB might have** a significant effect on performance
- ▣ **Carefully choose EPT** configuration: most Hypervisors uses a 2MB page EPT (internal fragmentation is better than memory access time degradation)
- ▣ Also **fine-tuning applications** require in many cases to use 2MB super-pages (or huge)
- ▣ **Constant improvements** in hardware generation to generation (and software, e.g. DBT, Unikernels)
- ▣ It's needed memory virtualization in **multi terabyte** RAM sytems? (e.g. GPT-3 requires ~1TB)



# How to address TLB pressure: Zen 4

