

# Secure Processor Architectures

---

## Chapter 3

- [1] J. Szefer, “Principles of secure processor architecture design,” *Synth. Lect. Comput. Archit.*, vol. 13, no. 3, pp. 1–173, 2018.

# Real-World Attacks

- ▣ Motivates the **need** for secure processor architectures
- ▣ Provide a **glimpse** of how wrong assumptions about hardware behavior (e.g., DRAM refresh) or unintended consequences of performance optimizations (e.g., speculative execution) affect security
- ▣ Same bugs and vulnerabilities of regular processors can affect secure processors too
  - ◆ Processor has bugs too, and sometimes are very expensive! (e.g., FDIV Pentium cost Intel millions)
  - ◆ Complexity implies bugs

# Real-World Attacks: Coldboot

- Used to stole information from RAM while the system is powered off
  - ◆ DRAM capacitors charge don't disappear suddenly when turned off: there is a slowly decay. **The assumption about DRAM fast volatility is wrong**
- If the DRAM chips are cooled (e.g., via compressed air) the decay is slower (capacitor decay is **temperature dependent**)
- Interchange the DRAM chips to another computer and dump the content of the chip it at **rogue OS boot**
  - ◆ **Stole keys, passwords, etc...**
- Solutions
  - ◆ Explicitly erase sensitive at power off. Use battery support to perform the operation
  - ◆ Encrypt memory content

# Real-World Attacks: Rowhammer

- ▣ Modify memory non-accessible from the attacker process (assuming OS/VM are ok)
- ▣ Bypass OS/VM isolation by exploiting DRAM cross/contaminations of row contents
  - ◆ A specific and **repetitive** access pattern to accessible memory to the attacker can “modify” adjacent non-accessible rows
- ▣ Identify the victim’s physical memory target (in an OS can be fixed) and try to allocate process data in an adjacent row (e.g., malloc across all memory). When achieved, **attack!**
  - ◆ **The attack might consist in code injection (write sniped of code that exploits system security).**
- ▣ Solutions:
  - ◆ Use better DRAM (hard do it. From generation to generation of DRAM the promises repeats, but the problem does not go away)
  - ◆ Prevent allocate the code in the same spots of the virtual-memory: **Address Space Layout Randomization** (ASLR)
  - ◆ Encrypt memory content

# Real-World Attacks: Meltdown

- ▣ Exploits **side effects of out-of-order execution** and design decisions of certain **Intel** processor families to read arbitrary memory of kernel (mapped in process address space)
  - ◆ Share kernel addressing space is advantageous for system calls
  - ◆ Share doesn't mean the user code can access (`pte.us=1`)
- ▣ Privilege level in memory access (loads) is only checked at commit stage (Intel processors)
- ▣ Execute loads to kernel addresses (stores the data in cache) and prevent the load to reach commit (e.g., raising an exception before to issue the load)
- ▣ Apply cache-side attack to the set of the cache that stores the data to infer the value
  - ◆ Piece-by-piece (~1 byte) using to timing determine the "stolen" address

# Real-World Attacks: Spectre

- ▣ Breaks isolation between apps by exploiting executive execution of instructions following branches
  - ◆ Like meltdown side-channel but affect any processor with branch-prediction
  - ◆ Harder to implement and cross application-attack
- ▣ Train in the attacker app the branch prediction (shared by all process running in the CPU) to force in the victim app a miss-prediction (branch prediction uses PC+history to make the prediction) using a sensitive piece of code (gadget)

```
if (x < array1_size)
    y = array2[array1[x]*256)
```


- ▣ Force the gadget to reach out-of-bounds access in array2 (will be cancelled later)
  - ◆ The remains of such access is in the cache and can be inferred via side-channel attack
- ▣ Solutions:
  - ◆ Disable branch predictor
  - ◆ Don't share branch predictor content between processes
  - ◆ Loads inside branch's acts as memory barriers

# Real-World Attacks: Other Bugs and Vulnerabilities

- Processor has bugs, documented by manufacturers. Some of them can be security-critical (30 on 300 in [4])
  - ◆ Examples System Management Mode, Message Signaling Interrupt, etc..
  - ◆ GPU vulnerabilities can be used also to break isolation and steal information
- Attacks can focus on non-compute components: Thermal Sensors, DVFS, can be abused. Change the timing of certain operations and introduce faults (to leak information)
- Many exploits design goals (performance, area and energy) are susceptible of being exploited
  - ◆ v.gr. Performance: caches via timing side-channel
  - ◆ v.gr. Area: DRAM via Rowhammer
  - ◆ v.gr. Power: DVFS can generate faults

# General-Purpose Architectures

- Secure processors (SP) are built on top of GP Processors (GPP) and expand them with security features
  - Its part of the Trusted Compute Base (**TCB**) Supports Trusted Execution Environment (**TEE**)
- GPP uses a ring-based protection mechanism to isolate App/OS/VMM
  - Certain ISA features available in each ring
  - When needed, controlled mechanism (syscalls, vmexits)
- Compromised OS or VMM can attack Apps.** SP tries to address that (untrusted OS/VM)



Green == Considered trusted



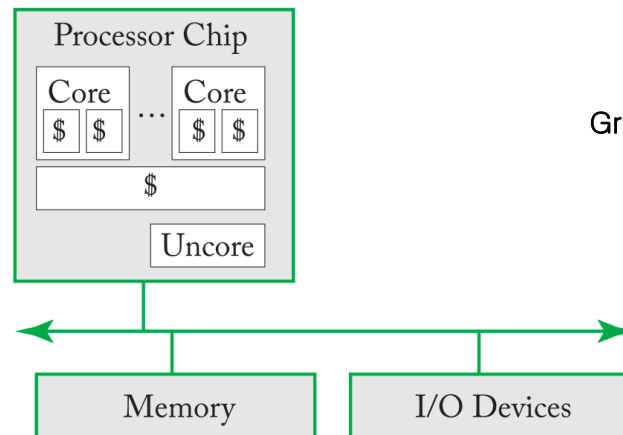
# Typical components (traditional view)

## ▣ Software Components

- ◆ Ring 1: Apps
- ◆ Ring 0: OS
- ◆ Ring -1: VMM

## ▣ Hardware Components

- ◆ CPU, Mem, complex I/O systems



Green == traditionally considered trusted


# Secure Processor Architectures (real view)

## ▣ **System Management Mode (SMM)** (ring -2)

- ◆ Code part of the firmware run by GPP
- ◆ Accessible via System Management Interface (SMI), asserting a pin in processor chip package or I/O over specific port
- ◆ Management functionalities (e.g., IPMI) even if OS/VMM compromised
- ◆ Uses security through obscurity

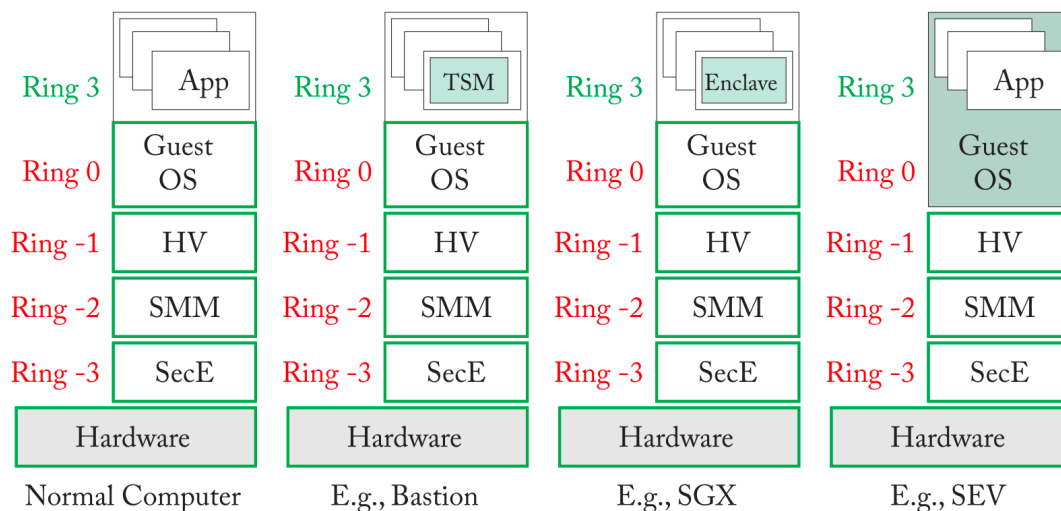
## ▣ **Platform Security Engine (SecE)** (ring -3)

- ◆ Intel Management Engine, AMD Platform Secure Processor (it's an ARM with TrustZone)
- ◆ **Small processor** isolated from the rest system (Intel in North-bridge), (AMD integrated)
- ◆ Can be online even with power-down
- ◆ Control system execution and emulate some hw features such as AMD SEV
- ◆ Uses security through obscurity



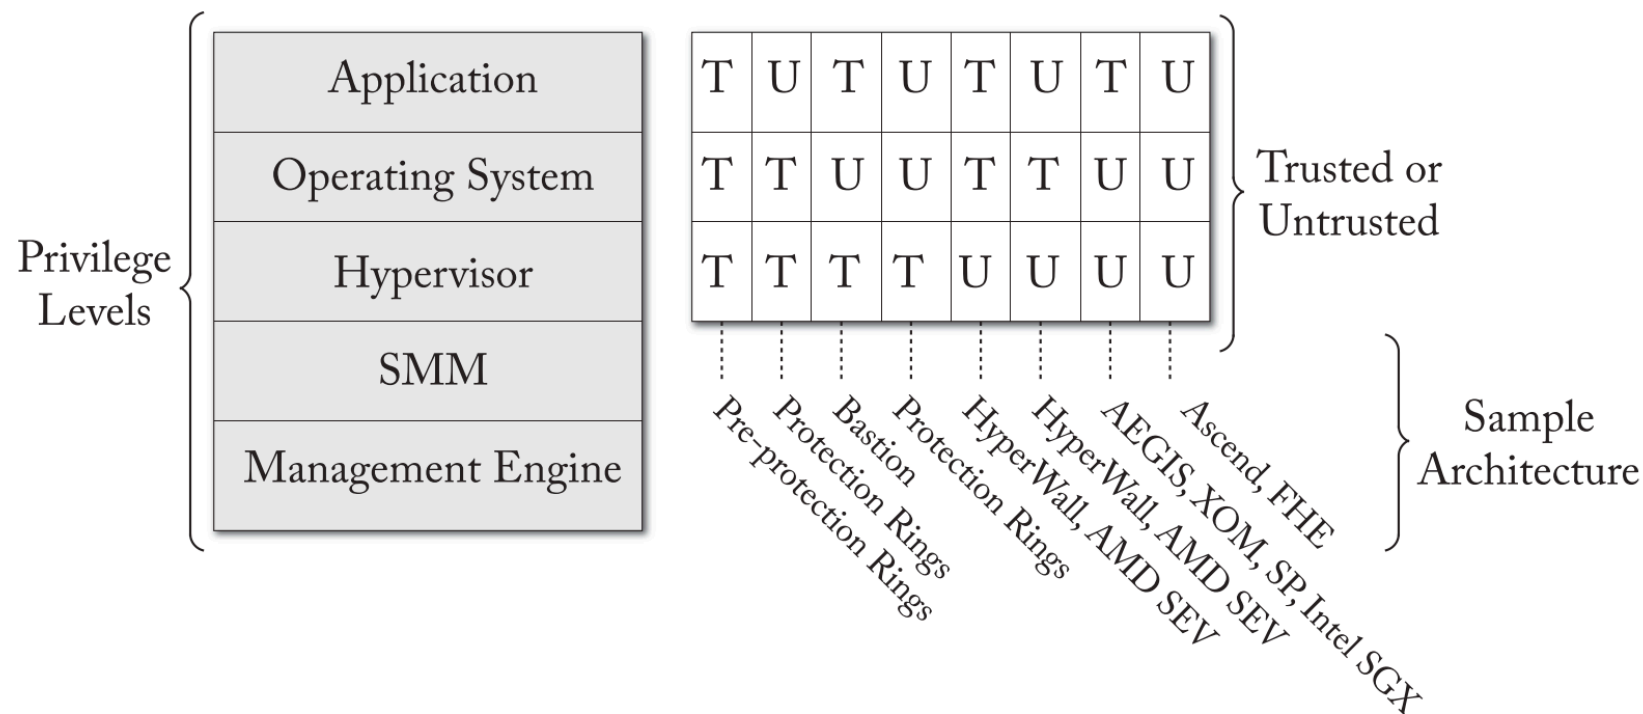
# Isolation Barriers

- SMM y SecE extends vertical privilege levels
- Horizontal privilege levels separation can be added also (e.g., ARM trustZone)
- Breaking vertical hierarchy of protection levels
- In secure mode (regardless of the level) software is more privileged than software in normal mode



# Architectures for Different Software Threats

- Diversifying the isolation, we can target specific attacks according to what is **trusted** or **untrusted** (assuming -2 and -3 is trusted)



# Architectures for Different Hardware Threats

## ▣ **Multiple chips** connected

- ◆ Susceptible of replacement or physical probing
- ◆ Some accessible (e.g., Memory) some don't (e.g., processor)
- ◆ In secure processor design memory and external wiring is untrusted
  - Modifications will be required to fix it

## ▣ **Processor is trusted**

- ◆ Too small to probe (5nm) with a reasonable cost
- ◆ An external bus is orders of magnitude easier to probe

## ▣ **3D and 2.5D** can make the system more resilient to hardware threats

- ◆ **Chip-let** designs (fewer external buses and chips in motherboard)
- ◆ Memory integration (e.g., Apple M1)

# Hardware Trusted Compute Base

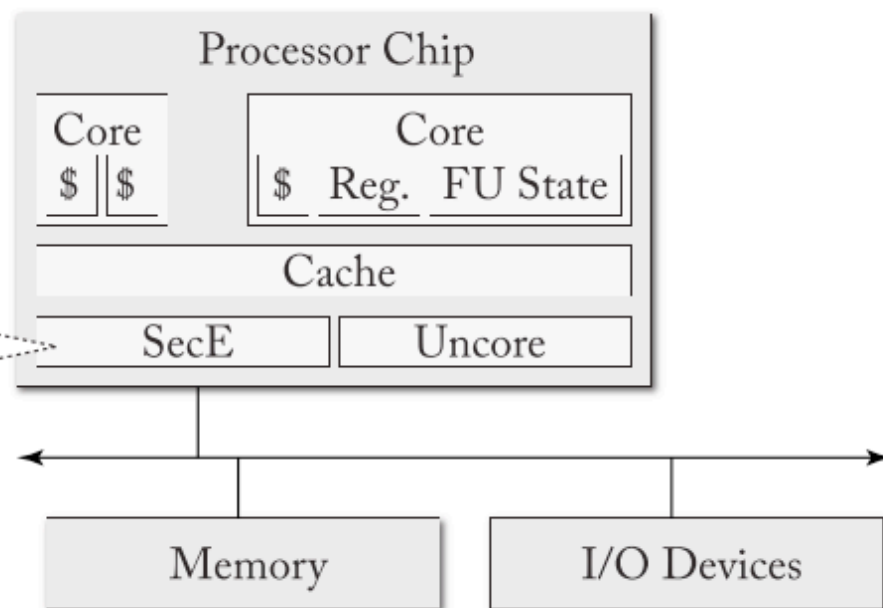
- As custom logic or dedicated processors

*Custom logic or hardware state machine:*

- Most academic proposals

*Code running on dedicated processor:*

- Intel ME = ARC processor or Intel Quark processor
- AMD PSP = ARM processor



# Examples of Secure Processor Architectures

## ▣ Academic

- ◆ XOM, AEGIS, NoHype, ...
- ◆ Initially targeting protecting software from hardware attacks (e.g., modification of off-chip memories)
- ◆ Protection against rogue OS added later and currently against rogue hypervisor too
- ◆ Some consider all system potentially rogue and compute without decrypting (e.g., homomorphic cryptography)
- ◆ Most focused in single-core systems

# Examples of Secure Processor Architectures

## ▣ Commercial

- ◆ 1970 IBM Logical Partitions
- ◆ Reconsidered in 2000s with IBM/Toshiba/Sony Cell Broadband Engine (PS3) (Security processor vault) and follows with ARM TrustZone, **Intel SGX**, **AMD SEV**, ...

- ▣ The pragmatic approach (processor) is flexible but also a weak point
  - ◆ The bugs in the software they run is vulnerable. The approach of security though obscurity amplifies the problem
  - ◆ Hardware solutions are too inflexible



# Secure Processor Architecture Assumptions

## ▣ Chip Assumptions

- ◆ It is the trust boundary for the hardware TCB
- ◆ Everything in the chip is trusted (and untrusted out of it)

## ▣ Size TCB Assumption

- ◆ Small software means less bugs, easy to verify and easier to audit
- ◆ Small Hardware, ""

## ▣ Open TCB Assumption

- ◆ Apply Kerckhoffs's Principle → no secrets int the TCB: has to be public.  
The only secret should be the cryptographic keys!

# Limitations of Secure Architectures

- ▣ Physical Realization Threats
  - ◆ Assume the manufacture is correct
  - ◆ Hardware Trojans might be added post-design in the foundry
  - ◆ Trojan detector can be included in the design too
- ▣ Supply Chain Threats
  - ◆ Current systems integrates many IP in the design/manufacture phase that can be integrated in late stages of the production
  - ◆ Use PUF (Physical Unclonable Functions) to verify that the system is compliant with the specification
- ▣ IP Protection and Reverse Engineering
  - ◆ Certain component might need to be "non-public"
  - ◆ Split-manufacturing (BEOL and FEOL in separate foundries)
- ▣ Side and cover attacks
  - ◆ Information leak trough unintended channels
- ▣ Alternatives to HW: **Full Homomorphic Encryption (FHE)**
  - ◆ Perform operation over cyphertext without leaking any information
  - ◆ Still not practical: currently very slow. Protects only the data:
  - ◆ If FHE is complete, TCB is no longer needed? There is no way to leak information with non-trusted hardware or software