

# Memory Protection

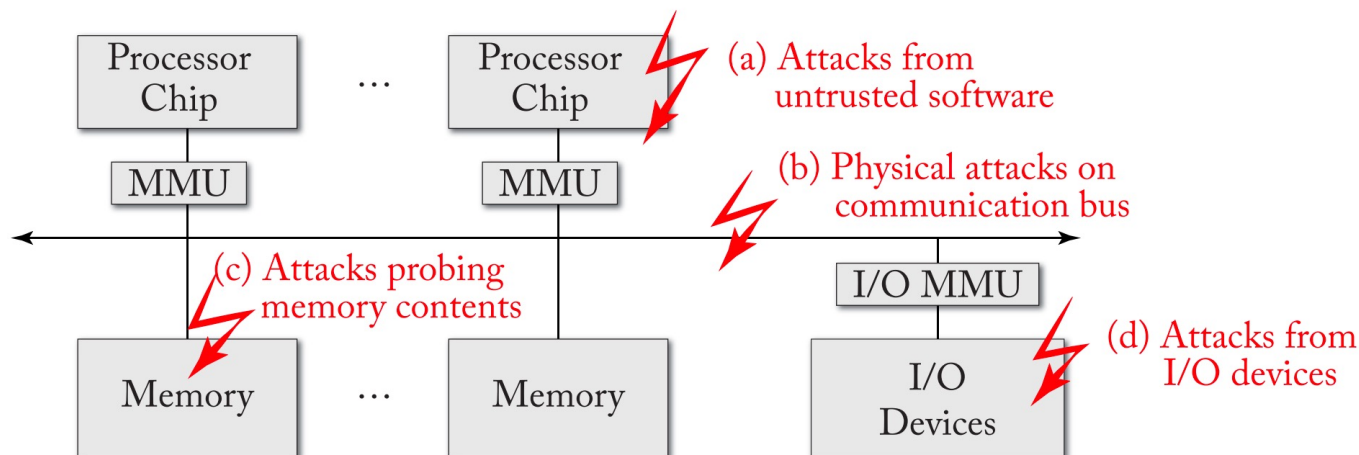
---

## Chapter 6

- [1] J. Szefer, “Principles of secure processor architecture design,” *Synth. Lect. Comput. Archit.*, vol. 13, no. 3, pp. 1–173, 2018.

# Threats Against Main Memory

- Wiring and memory device itself are untrusted assumed built upon DRAM devices
- Sources of Attacks
  - ◆ Untrusted software (bypassing isolation barriers)
  - ◆ Malicious devices (trying to access via DMA to protected regions)
  - ◆ Physical attacks (on memory "bus" via probing)
  - ◆ Physical attacks (on memory itself, Cooldboot, Rowhammer)
- Memory can be easily removed and analyzed off-line
- Competing Non-volatile Memories (NVM) in the horizon (such as Intel Optane) can make this problem harder



# Attacks to Memory

## ▣ Passive Attacks

- ◆ **Eavesdropping attacks:** observe information or accessing patterns without altering it to gain knowledge
  - Pattern: V.gr. Observing AES S-Box access pattern can leak information about the encryption key
  - Prevented: **encryption** and **obfuscation**

## ▣ Active Attacks

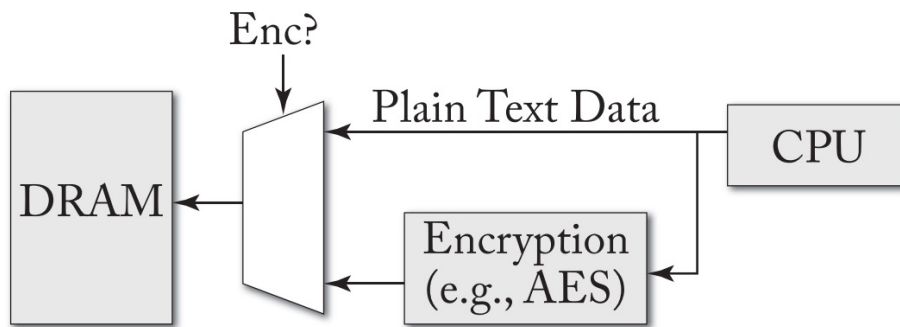
- ◆ **Spoofing attacks:** inject memory data (or operations) without being detected
  - Inject changes in memory mapping, processor instructions, stole data, ...
  - Prevented via **hashing**
- ◆ **Splicing Attacks:** combine multiple read/write operations in a new (legitimate) read or write
  - Splice parts of different messages (e.g., payload from one and header from other)
  - Prevented via **keyed hashing** (MAC Message Authenticated Code)
- ◆ **Replay Attacks:** Send messages again
  - Reuse old "known" messages (v.gr. replace an encryption key)
  - Prevented via **nonced hashing**

# Main Memory Protection Mechanisms

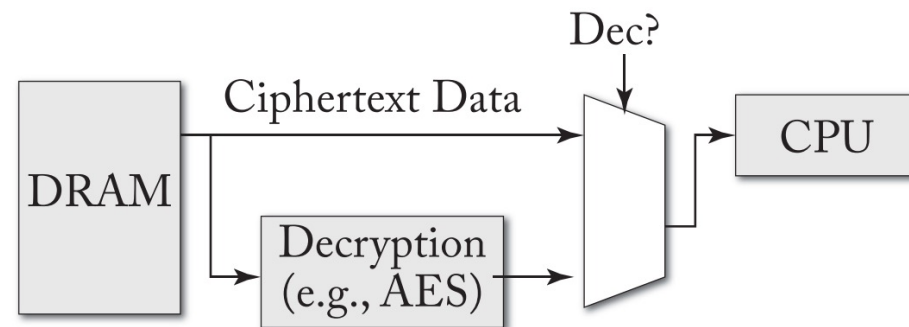
- ▣ Three main techniques and objectives
  - ◆ **Confidentiality**: with encryption
  - ◆ **Integrity**: with hashing (typically hash-trees)
  - ◆ **Access Pattern protection** with access pattern obfuscation
- ▣ Memory protection is focused on off-chip related issues
  - ◆ Encryption and/or hashing is only done at **off-chip processor interfaces**
  - ◆ Inside the processor chip regular isolation techniques (e.g., page tables for equally privileged software or tagging for higher privilege software)
  - ◆ Might increase average access time
  - ◆ Access pattern still discoverable
- ▣ MMU and IOMMU are also useful

# Confidentiality Protection With Encryption

- Use a **re-generated** key each time system **reboot**
- Store the key in the MMUs or special management module (usually Memory Controller)
- Application, OS or system level. Can be **selectively** enabled in some ranges of the addressing space of the process or physical addresses of the memory (use ASIDs in tags to isolate)
- Use asymmetric cryptography to support multi-socket systems (the key never should leave the processor chip in plaintext!!) or use independent key per memory controller/ASID
- Memory access will be slower for encrypted data (use **suitable algorithms** for the task: AES CTR Mode)
- To support DMA IOMMU also should have a key (I/O and PCIe is unencrypted, therefore vulnerable against probing)
- Can be vulnerable if management engine is untrusted (will have access to encryption keys)
- AMD Secure Encrypted Virtualization implements it



(a) Memory Write




(b) Memory Read

# Aside: AES Algorithm (10.000 feet view)

- Base on substitution-permutation network

- Example 3-rounds, 16-bit block



- AES

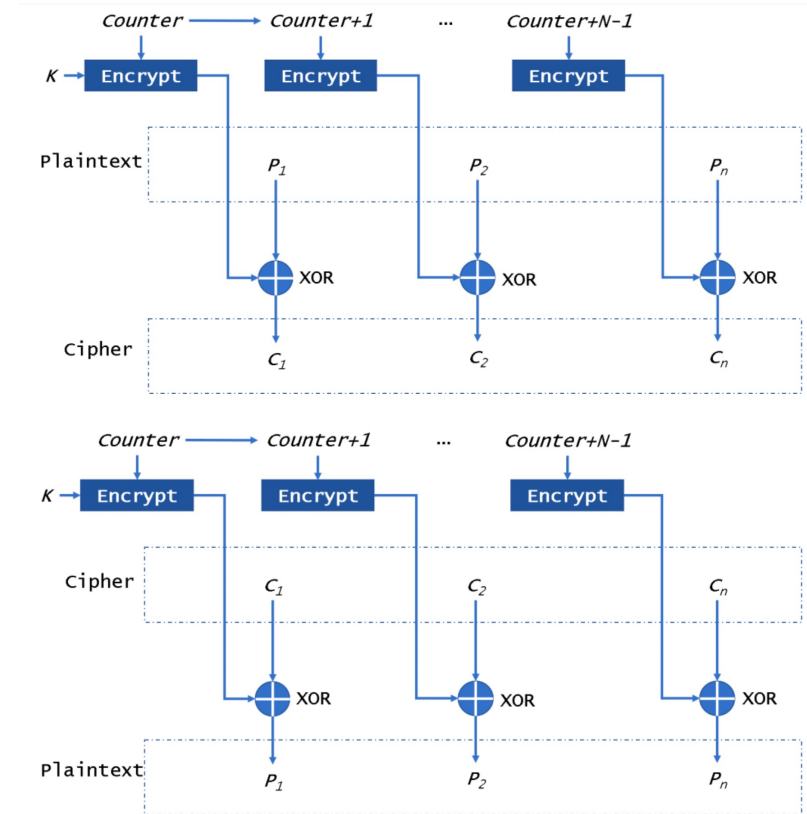
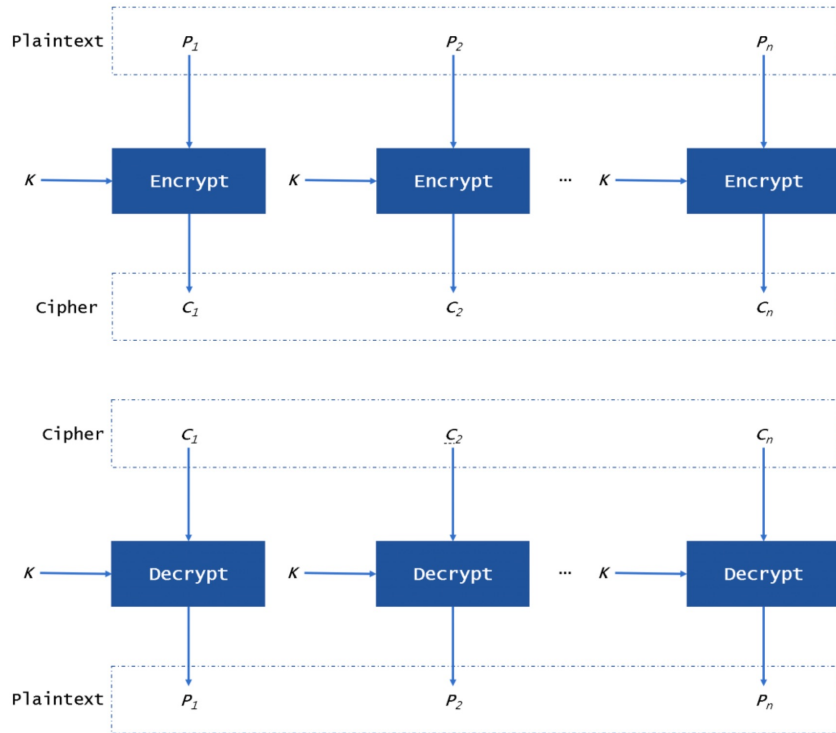
- 128-bit blocks
  - 128bit key: 10rounds
  - 256bit key: 14rounds
  - 192bit key: 12rounds
- Last round is special

- AES-NI

- Instructions for one round of ENC/DEC


# Aside: Counter Mode AES (fast encryption/decryption)

## ❑ Electronic Code Book vs Counter Mode (AES has 5 modes)



# Integrity Protection With Hashing

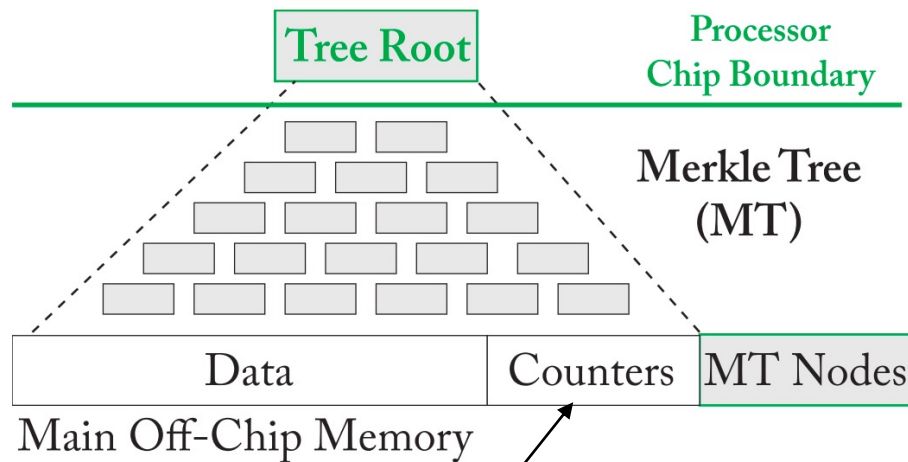
- ❑ Cryptographic hashes
  - ◆ Produces a unique output for an input (negligible probability of collision of two inputs if the hash is long enough)
  - ◆ We can't figure out the input from the output (non inversible functions)
- ❑ The result of hashing the data can be understood as a fingerprint of the data
  - ◆ The easiest way to check if the data has been altered is check hash result
- ❑ Integrity protections focused on **external attackers**
- ❑ Hash the whole memory is **impractical**: slow and every change requires to recompute the hash: **Use hash trees**
- ❑  $\log_2(N)$  hash computations (N: depth of the tree)
  - ◆ **Write**: update all the hashes from the leaves updated to the root
  - ◆ **Read**: check if involved hashes to the data are ok
- ❑ But attacker can modify the whole tree (??)
  - ◆ Add encryption to the hash





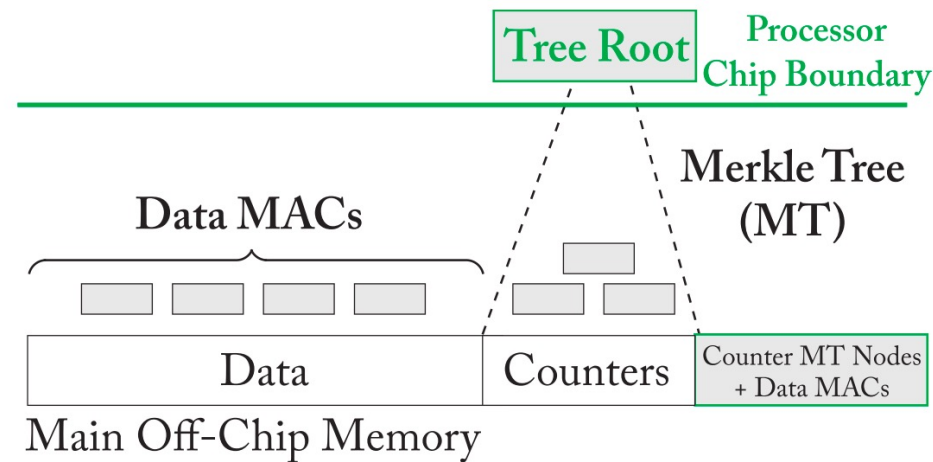
# Merkle Trees

- Root of tree on secure on-chip register



(a) Standard Merkle Tree

For AES CTR



(b) Bonsai Merkle Tree

# Hash Tree Protection

- ❑ Protect intermediate nodes in the tree: attacker can compute and insert its own hashes,
  - ◆ Do not use plain hashes: use secure processor-based keys to encrypt the hash (the tree is in memory!)
  - ◆ Use cryptographic hashes: **keyed hashes** (**MAC** or Message Authenticate Code) and always on secure on-chip locations: can't be tampered (but can be cached)
- ❑ MAC can protect against splicing and spoofing but **not replaying**
  - ◆ Add a monotonic counter for **replay attacks** protection
- ❑ **Secure page swapping** might require to reserve leaf nodes in the hash tree for non-present pages (i.e., integrity protection trees should cover the disk)
- ❑ Variety of **performance improvement** to minimize hashing performance impact
  - ◆ Counter mode encryption (out of the critical path)
  - ◆ Bonsais Merkle independent encryption seed from the memory address (do not combine key with address in the original tree)
- ❑ Intel SGX uses a tree of MACs for integrity and replay protection

# Access Pattern Protection

- Just observing the access to address activity can be used as conduit for revealing secrets
- ORAM (Oblivious RAM) (1984) keeps the semantics of the program but hides the access pattern by **shuffling memory locations** (inserting noise at compile time)
  - On each access memory insert random access. The objective is to make indistinguishable to the attacker the useful accesses from the noise
  - Only in sensible data (software + hardware, e.g. S-Box access in AES encryption)
- All OS does something called Address Space Layout Randomization (ASLR) but oriented to prevent exploits. Its at page level

