
Virtual Machines: Performance Optimization

Lab. 3

Outline

- Impact in performance and use of hardware assist
- Types of virtualizations with Xen

Task1 (20%)

- Create a minimal HVM installation disk domU with installed
 - Use a “lv” of at max 500MB
 - Configure the domU for “console” mode
- Configure the HVM machine to have access to PV installation (reuse /)

To “reuse” a disk image created in a HVM to boot a PV machine

- Copy hvm.cfg file to pv.cfg
 - change “kernel” and “ramdisk” by the existing Dom0 /boot
- Disk devices
 - Replace “xvda” by “hda” in “phy:”
 - Add root device to boot pv kernel (Assuming hvm installer used 1st partition)
 - root= /dev/xvda1 ro
- Fix console
 - Set extra options for pv kernel
 - extra="console=hvc0 "
 - Edit the /etc/inittab in image file and replace
 - 2:23:respawn:/sbin/getty 38400 tty2
 - By
 - 2:23:respawn:/sbin/getty 38400 hvc0
 - After that, both HVM and PV will work with a single image
 - **Remember do not boot simultaneously both VM!!**

The other way around: how to use a PV image on HVM

- Add a partition table to the PV volume
 - Many steps: parted, resizefs, grub, etc... (prone to errors)
 - Not recommended
- Our take: two volumes (one new volume with PT for /boot and use the PV for /)
 - Create /boot and install grub on it using the installer
 - Instruct the kernel to use PV volume on /
 - HVM will use the boot and PV is just the same!

Task 2: CPU Bound Workloads (30%)

- Compare the performance of two SpecCPU benchmarks selected in Lab1 using:
 - Native
 - PV without VT-X
 - PV with VT-x
 - HVM-on-PV
- Where the native performance went? (in the application with the strongest slowdown)
 - Just with PV with VT-x

Performance overheads analysis

- Unfortunately, PMU seems to be inaccessible to domU with current processors/Xen/kernels versions
 - `dmesg | grep -I pmu`
 - No driver found in kernel (?)
 - Little to no information on the webs...å
- Still the linux kernel has some interesting tracing support (and perf can access to it)
 - `perf list | grep xen`
 - xen:mc Hypercalls
 - Denoted as (mc) Multicalls due to xen “combine” multiple hypercalls in a single “multicall”
 - Key is “xen:xen_mc_entry”
 - xen.mmu Memory management
 - Includes shadow table, page allocation, page release, ...
 - Key are “xen:xen_mmu_flush_tlb , xen:xen_mmu_*_ptpage”
 - xen.cpu Context switches
- Kernel tracing is much more powerful and complex
- Caveat: With “HVM” (even with PV drivers) the kernel don’t expose those tracepoints
 - Requires the use of “xentrace”

Task 3: QoS for SpecJBB (20%)

- Configure two domU (using best virtualization) for running SpecJBB with latest Oracle JDK
 - Recompile the benchmark
- Configure CPU scheduling to achieve at least a 80% of native performance (SLA) in one of them (when executed concurrently)
 - Use remote ssh commands to synchronize execution
 - Make the configuration persistent
- (*) If answer is “impossible”...
 - How much is the minimum performance that we can guarantee?

Task 4: SpecWeb (30%)

- Compare the performance of SpecWeb2005
 - Native
 - HVM-on-PV
 - With separate domU for Clients+BeSim/Webserver
 - Add PCI-Pass-through for the network

- Multisystem run (PV-PCI-PT)
 - Configure besim and clients in a separate DomU
 - Run Dom0, and besim, clients in a single CPU (CPU0)
 - Migrate besim, client to another host (PCI-PT)

Grades

- Guide and presentation (25%)
- 75%
 - Content Lab2 (25%)
 - Content Lab3 (50%)